

A Dutch coreference resolution system with an evaluation on literary fiction

Andreas van Cranenburgh

University of Groningen, The Netherlands

November 14, 2019, Groningen, CL reading group

WHAT DO WE WANT?



**NATURAL LANGUAGE
PROCESSING**



WHEN DO WE WANT IT?



imgflip.com

WHEN DO WE WANT WHAT?



Plan for today

1. Background
2. Annotating Dutch novels
3. The coreference system
4. Evaluation
5. Future work

1. Background
2. Annotating Dutch novels
3. The coreference system
4. Evaluation
5. Future work

Definition

Coreference resolution is a structured prediction task of partitioning detected **mentions** into **entity clusters**.

Definition

Coreference resolution is a structured prediction task of partitioning detected mentions into entity clusters.

```
+-----+  
|           |  
"I voted for Obama because he was most aligned with my values", she said.  
|                                           |   |  
+-----+
```

Definition

Coreference resolution is a structured prediction task of partitioning detected mentions into entity clusters.

[illegible]

- ▶ Entity 1 = {Obama, he}
- ▶ Entity 2 = {I, my, she}

Mentions

Definition

Mention or referring expression: span of text that refers to a person or object in the real or mention world.

NB: contrast with **markable**, a *potentially* referring expression.

Mentions

Definition

Mention or referring expression: span of text that refers to a person or object in the real or mention world.

NB: contrast with **markable**, a *potentially* referring expression.

Pronouns I, he, my, his, that, [each other], himself, ...

Names [John], [John Smith], [Mr. Smith], ...

Nominals [the man], [the flowers on [the table]], ...

But not:

- ▶ Events, actions, times
- ▶ Non-referring expressions: it rains

Coreference relations

Strict [John]₁ sees [Mary]₂.
[He]₁ waves at [her]₂ and takes [[his]₁ bike]₃.

Predicative [John]₁ is [the director]₁.

Appositions [John]₁, [the painter]₁. [Bond]₂, [James Bond]₂.

Relative clauses [The man]₁ [who]₁ sold [the world]₂ ...

Bound anaphora [Everyone]₁ has [his]₁ opinion about it.

Coreference relations

Strict [John]₁ sees [Mary]₂.
[He]₁ waves at [her]₂ and takes [[his]₁ bike]₃.

Predicative [John]₁ is [the director]₁.

Appositions [John]₁, [the painter]₁. [Bond]₂, [James Bond]₂.

Relative clauses [The man]₁ [who]₁ sold [the world]₂ ...

Bound anaphora [Everyone]₁ has [his]₁ opinion about it.

- ▶ No bridging relations (e.g., part-whole)
- ▶ Relation type not annotated

Winograd schemes

The [city councilmen]₁ refused [the demonstrators]₂ a permit because ...

1. ... [they]₁ feared violence.
2. ... [they]₂ advocated violence.

“AI-complete” problem

History of coreference resolution

Various datasets, languages:

1996 MUC-6 shared task, English

2004 ACE shared task, English/Chinese/Arabic

2010 SemEval shared task, multilingual including Dutch

2011 CoNLL shared task, English

2012 CoNLL shared task, English/Chinese/Arabic

State of the art: from rules to a neural arms race ...

Data: OntoNotes (English)

CoNLL score

CoNLL 2011	shared task, winner: Lee et al., rule-based	58.3%
CoNLL 2012	shared task, winner: Fernandes et al., perceptron	58.7%
EMNLP 2017	end-to-end coref. resolution, deep learning	67.2%
NAACL 2018	e2e + ELMO + c2f, deeper learning	73.0%
EMNLP 2019	e2e + BERT Large, even deeper learning	76.9%

Evaluation metrics

Coreference evaluation is a **mess**!

Fatally flawed metrics:

1996 MUC

1998 B³

2005 CEAFm, CEAF_e

2011 CoNLL score (= avg of MUC, B³, CEAF_e)

2011 BLANC

No known issues (yet!):

2016 Link-based Entity-Aware metric (LEA)

Moosavi & Strube (ACL 2016) Which coreference evaluation metric do you trust?
A proposal for a link-based entity aware metric

1. Background
2. Annotating Dutch novels
3. The coreference system
4. Evaluation
5. Future work

By the way ...

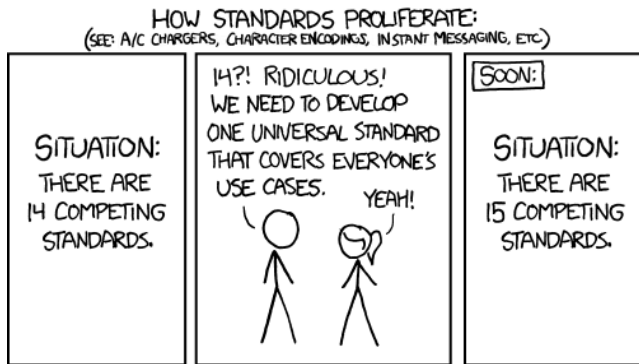
#BenderRule:

The rest of this talk is about Dutch!

<https://thegradient.pub/the-benderrule-on-naming-the-languages-we-study-and-why-it-matters/>

I made a new annotation scheme ...

I made a new annotation scheme ...



<https://xkcd.com/927/>

Annotation scheme

Simplified annotation scheme:

- ▶ Annotate mentions: include singletons, exclude non-referring expressions.
- ▶ Avoid difficult mention boundaries: no discontinuity, relative clauses
- ▶ Only annotate entity clusters, not directed anaphor-antecedent relations

Annotation workflow

1. Tokenize, parse with Alpino
2. Run coreference system
3. Manually correct output with CorefAnnotator
4. Optional: correction by second annotator

<http://www.let.rug.nl/vannoord/alp/Alpino/>
<https://github.com/nilsreiter/CorefAnnotator/>

Annotation workflow

1. Tokenize, parse with Alpino
2. Run coreference system
3. Manually correct output with CorefAnnotator
4. Optional: correction by second annotator

Result: tabular CoNLL 2012 file

```
#begin document (example); part 000
example 1      John      (0)
example 2      sees      -
example 3      Mary      (1)

#end document
```

Annotated texts

	CLIN26 dev set	SemEval 2010 dev	Novels, dev set	Novels, test set
documents	30	23	10	11
tokens	4018	9164	19,051	88,092
sents per doc	7	21.4	100	491.5
avg sent len	19.3	18.4	19.0	16.3

Annotated texts

	CLIN26 dev set	SemEval 2010 dev	Novels, dev set	Novels, test set
documents	30	23	10	11
tokens	4018	9164	19,051	88,092
sents per doc	7	21.4	100	491.5
avg sent len	19.3	18.4	19.0	16.3
entities	273	424	1798	8337
mentions	663	1010	4243	20,873
% pronouns	7.69	14.45	43.3	36.5
% nominal	52.34	54.35	46.2	52.2
% names	39.97	31.20	10.5	11.2

107k tokens of annotated literary text!

1. Background
2. Annotating Dutch novels
3. The coreference system
4. Evaluation
5. Future work

Mention detection

1. Extract candidate constituents
 2. Adjust spans
 3. Filter with patterns
 4. Detect features
- ▶ [John], [who], [the book on [the table]], ...
 - ▶ [Mr Smith], [San Jose, California], ...
 - ▶ Drop: yesterday, about two meters, ...
 - ▶ gender, animacy, number

Mention feature detection

parse features in the Alpino parse tree (HPSG-inspired)

NER part of Alpino; person/org/loc/misc

wordnet animacy & gender of head nouns (hand-corrected)

web text names extracted w/heuristics from 30GB parsed **English** text
(Bergsma & Lin 2006; TODO: replicate for Dutch!)

Mention feature detection

parse features in the Alpino parse tree (HPSG-inspired)

NER part of Alpino; person/org/loc/misc

wordnet animacy & gender of head nouns (hand-corrected)

web text names extracted w/heuristics from 30GB parsed **English** text
(Bergsma & Lin 2006; TODO: replicate for Dutch!)

	Pronouns	Nominals	Names
Number	parse	parse	web text
Gender	parse	wordnet	web text
Animacy	parse	wordnet	NER, web text

Sieves: link mentions with deterministic rules

Quote attribution find speaker & addressee of direct speech

String match [The boy]₁ ... [The boy]₁ ...

Precise constructs [the boy]₁ [who]₁ ...

Head match [The clever boy]₁ ... [the boy]₁

Proper head noun match [Bond]₂, [James Bond]₂

Pronoun resolution [He]₁ ... [his]₁ ...

Demo

file:///tmp/Voskuil.html

Move mouse over bracketed text to highlight coreference. Move mouse over direct speech to highlight speaker and addressee. Click on a sentence to toggle the display of its parse tree.
Legend: [Singleton] [Coreference] [Speaker] [Addressee] ' Direct speech '

TOP
DU
SV1-nuc1
PPART-vc
SV1-tag
let ww-hd vnw-su vnw-obj1 bw-mod ww-hd let let ww-hd vnw-su let
Heb je hem dan gezien ? vroeg ik .

[XML source](#)

In [het achterhuis] was [een groothandel in wc-potten] gevestigd . Er werkte één man . [Hij] kwam om negen uur , als [ik] al naar [[mijn] werk] was , en vertrok om vijf uur , voor [ik] terugkeerde . [Nicolien] hoorde [hem] langskomen als [ze] bezig was met [de afwas] . [Hij] kwam dan over [het portaal] , klom de negen treden naar [het achterhuis] op , opende [[zijn] voordeur] en sloot [haar] zachtjes achter [zich] . De rest van de dag merkte [ze] niets van [hem] , tot [hij] weer wegging . Er kwamen ook geen bezoekers .

' Het is [een oude man] , denk [ik] , ' zei [ze] .

' Heb [je] [heny] dan gezien ? ' vroeg [ik] .

' Nee , dat kan ...

8 3 human=1 gender=m number=sg person=3 inquote=1 necclass=None
head=hem
Er werd gebeld ... , in een
slecht zittend , wat te ruim , grjs kostuum .

<https://andreasvc.github.io/voskuil.html>
<http://haytabo.let.rug.nl:8888/> (RuG only)

1. Background
2. Annotating Dutch novels
3. The coreference system
4. Evaluation
5. Future work

Evaluation: shared tasks

CLIN26 shared task	Mentions	BLANC			
GroRef, Boeing test set	59.34	30.96			
This Work, Boeing test set	59.49	31.48			
GroRef, GM test set	60.40	31.31			
This Work, GM test set	59.26	31.07			
GroRef, Stock test set	53.70	25.40			
This Work, Stock test set	54.68	26.09			
SemEval 2010, Dutch, test set	Mentions	BLANC	MUC	B ³	CEAFm
SemEval 2010: Sucre	42.3	46.9	29.7	11.7	15.9
SemEval 2010: UBIU	34.7	32.3	8.3	17.0	17.0
This Work	64.27	41.48	51.95	45.85	51.20

Evaluation: novels

	mention F1	recall	precision	LEA F1
SemEval 2010 (test set)	64.27	36.00	39.96	37.88
CLIN26 shared task (Boeing test set)	59.49	29.83	33.95	31.76
Literary texts (dev set)	87.05	57.13	61.71	59.33
Literary texts (test set)	87.10	49.27	57.45	53.05

Discussion

- ▶ High variance among novels; matter of style?
- ▶ Better performance on novels than news! Surprising?
 - ▶ More dialogue and pronouns in novels (some long chains)
 - ▶ Novels are longer documents (including our annotated fragments)
 - ▶ Not all errors are created equal ...

1. Background
2. Annotating Dutch novels
3. The coreference system
4. Evaluation
5. Future work

Improve components

Components:

- ▶ Mention detection/spans
- ▶ Pleonastic pronoun detection
- ▶ Gender/animacy
- ▶ Quote attribution
- ▶ Pronoun resolution

Improve components

Components:

- ▶ Mention detection/spans
- ▶ Pleonastic pronoun detection
- ▶ Gender/animacy
- ▶ Quote attribution
- ▶ Pronoun resolution

Procedure:

1. Acquire/annotate more data
2. Train supervised classifier
3. ???
4. Profit!

Neural coreference

- ▶ Create/convert more annotated data
- ▶ Train End-to-end coreference system
- ▶ with BERT ...

Neural coreference


- ▶ Create/convert more annotated data
- ▶ Train End-to-end coreference system
- ▶ with BERT ...

Downsides:

- ▶ Huge memory (32GB GPU) & carbon footprint
- ▶ Long documents need to be split into paragraph chunks ...

Kenton Lee et al. (EMNLP 2017) End-to-end neural coreference resolution

RULE-BASED



STATISTICAL



NEURAL



BERT



THE END

Code: <https://github.com/andreascv/dutchcoref>

Paper: Coming soon™

Thanks to my BSc thesis students for helping with annotation!

